

Statistical Science.

Part III. The General Linear Model.

Chapter 9.1 Regression. Explanatory Variable Fixed by Experiment

ReCap.	Part I (Chapters 1,2,3,4)
ReCap	Part II (Ch 5, 6, 7)
ReCap	Part III (Ch 8 GLM components)
9.1	Explanatory Variable Fixed by Experiment
9.2	Explanatory Variable Fixed into Classes
9.3	Explanatory Variable Measured with Error
9.4	Exponential Functions
9.5	Power Laws. Linear Regression
9.6	Model Revision

Data files & analysis SC_9_3_1.out Ch9.xls
--

on chalk board

ReCap Part I (Chapters 1,2,3,4)

Quantitative reasoning: Example of scallops,
which combined models (what is the relation of scallop density to substrate?)
with statistics (how certain can we be?)

ReCap Part II (Chapters 5,6,7)

Data equations summarize pattern in data as a series of parameters (means, slopes).
Frequency distributions, a key concept in statistics, are used to quantify uncertainty.
Hypothesis testing uses the logic of the null hypothesis to make a decision about an
unknown population parameter.

Estimation is concerned with the specific value of an unknown population parameter.

Today: The General Linear Model Regression: Single explanatory variable fixed by experiment. Work through this example, using a generic recipe.

Wrap-up

Regression a special case of the general linear model.

Response variable as a function of a single explanatory variable.

Relation between variables is expressed as a slope.

H_A/H_0 pair about the slope parameter.

The variance in the response variable is partitioned according to the model.

The partitioning is used to evaluate the fit of the model to the data.

The inverse of the unexplained variance is used to calculate the likelihood ratio.

The ratio of explained to unexplained variance is used to calculate the F-ratio.

The F-ratio is used to calculate the p-value.

This is a measure of uncertainty on the measure of evidence, the likelihood ratio.

This measure of uncertainty is used to declare a decision about the null hypothesis.

GLM regression. Example: An experimental study with single fixed variable.

Example 9.3.1 from Snedecor and Cochran (1989). The quantity of interest is the phosphorus content of corn (P_{corn} in ppm), in relation to the phosphorus levels in samples of soils with experimentally fixed levels of phosphorus (P_{soil} in ppm).

Does the phosphorus content of corn increase when soil phosphorus is increased ?

P_{soil}	P_{corn}
1	64
4	71
5	54
9	81
13	93
11	76
23	77
23	95

1. Construct model

Constructing a model is unfamiliar to most students so we do it step by step. We will start with a verbal model, proceed to graphical, and from there write the formal model.

We begin by listing the variables with a symbol and type of units.

P_{corn} = Phosphorus content of corn (ppm).

Continuous variable on a ratio type of scale.

P_{soil} = phosphorus content of prepared sample of soil (ppm).

Continuous variable on a ratio type of scale.


State the verbal model: Phosphorus content of corn depends on phosphorus content of soil.

Sketch a graphical model: A straight line relation.

The sketch distinguishes the response (Y-axis) from the explanatory variable (the X axis).

Next, table the variables by role

Name	Units	Type	Role
P_{corn}	ppm	Ratio	Response
P_{soil}	ppm	Ratio	Explanatory



The response variable is called the dependent variable.

The explanatory variable is called an independent variable.

This defines the mathematical relation.

This can be confusing because “independent” variables can depend on each other.

So response/explanatory is preferable. It arises from the scientific question.

Separating the response from explanatory variables is the first and most important step in statistical analysis. If someone comes to you for statistical advice, the best way to help is often to ask them to state the response variables, and to separate these from the explanatory variables. This clears away the fog that surrounds the search for the “right statistical test.”

1. Construct model

Use the table of variables to construct the model.

		<u>Units</u>	<u>Dimensions</u>	<u>Type of measurement scale</u>
Response	P_{corn}	(ppm)	dimensionless	ratio
Explanatory	P_{soil}	(ppm)	dimensionless	ratio

The scale type of the response variable determines the analytic procedure.	
Response variable	
if nominal	→ frequencies, then Generalized Linear Model GzLM
if nominal	then median test, binomial tests, etc
if ordinal	then "non-parametric" tests. e.g. Kruskal-Wallis
if ratio	then GLM
The scale of the explanatory variables determine type of model.	
Explanatory variable in GLM (normal error)	
if nominal	then ANOVA
if ordinal	then ANOVA
if ratio	then regression
if ratio and otherwise	then ANCOVA
Explanatory variable in GzLM (normal and non-normal errors)	
if nominal	then ANODEV, G-tests, and extensions
if ordinal	then ANODEV, G-tests and extensions
if ratio	then logistic regression
if ratio and otherwise	then ANODEV, likelihood ratio tests

Now write the model using names of quantities.

$$P_{corn} = f(P_{soil}).$$

"Phosphorus in soil depends on phosphorus in the soil"

Finally, write the model in more abstract form, which is what the computer will use to carry out the analysis.

$$P_{corn} = \alpha + \beta_{P_{soil}} \cdot P_{soil} + \varepsilon$$

With practice, this sequence becomes practically automatic.

The model is for the population. α and $\beta_{P_{soil}}$ are the parameters for the population.

α is the Y-axis intercept.

$\beta_{P_{soil}}$ is the slope parameter.

	<u>Units</u>	<u>Dimensions</u>	<u>Type of measurement scale</u>
α	same as P_{corn}	none (mass/mass)	ratio
$\beta_{P_{soil}}$	(ppm/ppm)	none (ppm/ppm)	ratio

1. Construct model

Here is the model for the sample. It uses parameter estimates from data.

$$P_{corn} = \hat{\alpha} + \hat{\beta}_{P_{soil}} \cdot P_{soil} + residual$$

$\hat{\alpha}$ and $\hat{\beta}_{P_{soil}}$ are estimates of the parameters α and $\beta_{P_{soil}}$

An alternative convention is to use greek letters for population parameters, roman letters for estimates. Using this convention, the model for the sample is written with roman letters:

$$\begin{array}{ll} P_{corn} = \alpha + \beta_{P_{soil}} \cdot P_{soil} + \varepsilon & \text{Population} \\ P_{corn} = a + b_{P_{soil}} \cdot P_{soil} + residual & \text{Estimates from sample} \end{array}$$

This convention falls apart after 2 letters. α, β, γ to a, b, g ?

So we will use hats on top of the greek symbols for estimates of these parameters.

2. Execute analysis. Place data in model format:

Column of data for response variable labelled P_{corn}

Column of data for explanatory variable labelled P_{soil}

Code model statement in statistical package according to the GLM

$$P_{corn} = \alpha + \beta_{P_{soil}} \cdot P_{soil} + \varepsilon$$

```
MTB> regress 'Pcorn' 1 'Psoil'
MTB> GLM Pcorn = Psoil.
```

All packages use a model statement to code the GLM. In some packages this model statement is typed (SAS, Minitab, R). In other packages it is present but not obvious. The example in the box shows the coding in Minitab for two different commands that produce the same result. Some packages have a graphics interface that allows you to code this model (e.g., SPSS). If you are using the graphics interface, it helps to look at the code produced, so that you understand how the model you produce with the graphic interface translates into a model statement.

2. Execute analysis. Compute fitted values and residuals.

The model statement is similar across statistical packages.

Extracting the residuals and fits differs among packages.

```
MTB > GLM Pcorn = Psoil;
SUBC> fits fits;
SUBC> res residuals.      Minitab
```

```
Proc GLM;
  Model Pcorn = Psoil      SAS
```

```
CornModel <- lm(Pcorn ~ Psoil)
                                                R
```

```
Regress Pcorn Psoil,
                                                Stata
```

```
MTB > print 'Pcorn' 'fits' 'residuals'
ROW  Pcorn    fits    res
1     64    62.997    1.0031
2     71    67.248    3.7524
3     54    68.665   -14.6645
4     81    74.332    6.6679
5     93    80.000   13.0003
6     76    77.166    1.1659
7     77    94.169   -17.1687
8     95    94.169    0.8313
9    109   101.253    7.7468
```

2. Execute analysis. Compute fitted values and residuals.

Here is the computational sequence:

1. Estimate the parameters
2. Calculate fitted values from the parameter estimates, for each data equation.
3. Calculate the residuals (response variable – fitted value).

These are readily obtained in from model based routines in statistical packages. We depend on the package to make these calculations correctly. Here is a brief tour of the machinery, for those who are interested.

1. Estimate parameters α and β_{Psoil} from the sample. The least squares estimate of β_{Psoil} minimizes the sum of the squared residuals (deviations of the data from the line).

$$SSE = \sum_{i=1}^n (Y - \hat{Y})^2 = \sum_{i=1}^n (Y - \hat{\alpha} - \hat{\beta}X)^2$$

where there are n observations indexed by i . Some routines use iterative search: make a guess, compute the SS_{error} , make another guess, compute SS_{error} for this guess, compare to previous SS_{error} , continue until SS_{error} is as small as possible.

For a simple straight line model most routines obtain the estimate of $\hat{\beta}$ from the following formula.

$$\hat{\beta} = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^n (X_i - \bar{X})^2}$$

For the Corn data the estimate is $\hat{\beta}_{Psoil} = 1.4169$

To estimate the y-intercept α we use the mean values of the response and explanatory variable.

$$\text{mean}(Pcorn) = \hat{\beta}_0 = 80 \text{ ppm}$$

$$\text{mean}(Psoil) = 13.0 \text{ ppm}$$

$$\hat{\alpha} = \hat{\beta}_0 - \hat{\beta}_{Psoil} (\text{mean}(Psoil)) = 80 - 1.42(13)$$

$$\hat{\alpha} = 61.58 \text{ ppm}$$

Here is the relation of the regression equation (which has the y-intercept α) to the GLM (which has β_0 the mean value of the response variable).

$$\text{GLM: } Pcorn - 80.0 = 1.42 (Psoil - 13.0) + \text{res}$$

$$\text{Regression Eq: } Pcorn = 61.58 + 1.42 Psoil + \text{res}$$

2. Calculate the fitted values from the parameter estimates,

$$\text{Fitted values: } \text{Fits} = E[Pcorn] = \hat{\alpha} + \hat{\beta}_{Psoil} \cdot Psoil$$

3. Calculate the residuals from the fitted values:

$$\text{Residuals: } \text{Res} = Pcorn - \text{Fits}$$


3. Evaluate Model.

We use the residuals to evaluate the model.

We begin by examining the straight line assumption. Is this valid?

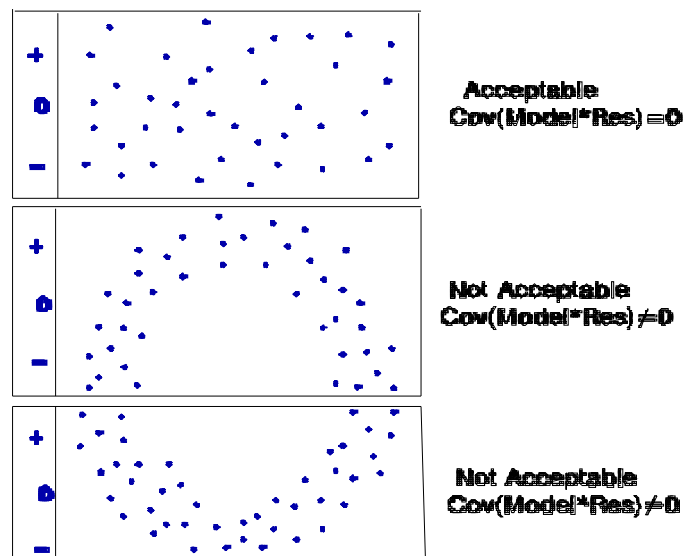
We use the residual versus fit plot to evaluate this assumption.

If the assumption is valid, the plot will show a band from left to right. If the assumption is not valid, the plot will show either a bowl or an arch pattern. Looking at residual fit plot we see that the straight line model is acceptable for the corn data.



Here is a diagram that contrasts an acceptable plot with two unacceptable plots. The acceptable plot shows a uniform band. The unacceptable plots show either an arch or a bowl. Bowls or arches result if the relation of the response to explanatory variable is curvilinear.

If this assumption is violated we go back to step 1 and reformulate the model to something other than a straight line.



3. Evaluate model.


Next, we evaluate the error model that was used to estimate the parameters and likelihood ratios that will be used to calculate p-values (Type I error) from a statistical distribution (chisquare or t or F).

These distributions assume:


- Fixed variance (errors homogeneous)
- Normally distributed errors.
- Independent errors
- Unbiased estimate (errors sum to zero)

In practice the first assumption is the most important. The second assumption is more often mentioned and often incorrectly diagnosed (*i.e.* diagnosed before estimating parameters and computing residuals). The third assumption is best diagnosed if we know the order in which samples were gathered or we know the spatial arrangement of samples. The fourth assumption does not need to be checked when parameters are estimated by statistical packages that automatically produce unbiased estimates. We will focus on the first two assumptions, unless we have information allowing us to diagnose the independent error assumption.

To evaluate the fixed error assumption we again examine the residual versus fit plot. If the assumption is valid the plot will show a horizontal band. The dispersion around zero will be uniform across the plot. If the assumption is not valid the plot will show vertical dispersion that changes from left to right in the plot, usually with an obvious cone, fan, or spindle pattern. The residual fit plot for the corn data shows an acceptable band, with no evidence of change in dispersion going from left to right. The assumption of homogeneous error is acceptable for the corn data.



Here is a diagram shows acceptably homogenous residuals in idealized form, superimposed on the corn data. The dispersion around the regression line is equal all along the line. As a result, the residual versus fit plot shows a uniform band running from left to right.



This diagram shows a pattern of heterogeneity in which the dispersion increases from left to right along the regression line. This pattern, which is common, will result in a cone opening out to the right in the plot of residuals versus fitted values. The opposite pattern, of a cone opening out to the left due to high variance at small fitted value, is rare.

3. Evaluate model. Normal Errors


The next assumption is that the errors are normal. The histogram of residuals for the corn data looks “somewhat” normal, but there is evidence of skew to low values.

Another diagnostic is comparison of the observed distribution to a normal distribution with the same mean and standard error.

```
MTB > hist 'res'
```

Histogram of res N = 9

Midpoint	Count	
-15	2	**
-10	0	
-5	0	
0	3	***
5	2	**
10	1	*
15	1	*



A useful plot for diagnosing normality shows each residual plotted on a normal probability scale. If the residuals are normal, the residuals will fall on a straight line rising from left to right. The normal probability plot for the *Pcorn* analysis shows two residuals that deviate from the diagonal. These fall within the bounds of the first (25%) and third (75%) quantiles.

3. Evaluate error model. Independent Errors

Because this is a text example, we do not have information on spatial layout of samples, or on collection sequence. We will assume independent errors. Just because our observations come from a single location, we do not assume that our data are ‘not independent.’ See Hurlbert (1984) for an example of this fallacy.

3. Evaluate error model. Conclusion.

Residuals appear to be homogeneous with acceptable departure from normal distribution. We do not have enough information to evaluate the assumption of independent errors. Based on graphical evaluation we judge the residuals acceptable. To check this judgement we will compare the p-value assuming normal error to the randomized p-value (no assumption of normal error).

4. Partition df and SS according to model.

Compute total degrees of freedom

$$df_{\text{total}} = n - 1 = 9 - 1 = 8$$

Partition df_{total} according to model, using rules

$$df_{\text{model}} = 1 \quad \text{regression line}$$

$$df_{\text{res}} = df_{\text{total}} - df_{\text{model}} \quad df_{\text{res}} = 8 - 1 = 7$$

Model	$P_{\text{corn}} - \beta_0$	=	$\beta_{P_{\text{soil}}} \cdot P_{\text{soil}}$	+	ε
Source	Total	=	P_{soil}	+	Resid
df	$9 - 1$	=	1	+	7

General linear model routines estimate the total sum of squares, partitioned according to the model.

GLM	$P_{\text{corn}} - \beta_0$	=	$\beta_{P_{\text{soil}}} \cdot P_{\text{soil}}$	+	ε
Source	Total	=	P_{soil}	+	Resid
df	$9 - 1$	=	1	+	7
SS	2274	=	1473	+	800.4
	SS_{tot}	=	SS_{regr}	+	SS_{res}

Here are calculations from data equations.

Data Equations for null model, $P_{\text{corn}} = \text{mean}(P_{\text{corn}})$

	Data =	Model	+ Res	Res ²	
	64.00	80.00	-16.00	256.0000	
	71.00	80.00	-9.00	81.0000	
	54.00	80.00	-26.00	676.0000	
	81.00	80.00	1.00	1.0000	
	93.00	80.00	13.00	169.0000	
	76.00	80.00	-4.00	16.0000	
	77.00	80.00	-3.00	9.0000	
	95.00	80.00	15.00	225.0000	
	109.00	80.00	29.00	841.0000	
Sums	720.00	720.00	0.00	2274.0000	= sum(res ²)
	720/9 = 80				

Data Equations for regression model

$P_{\text{corn}} = 61.58 + 1.417 \cdot P_{\text{soil}}$

	Psoil	Data =	Model	+ Res	Res ²		
	1	64.00	63.00	1.00	1.0055		
	4	71.00	67.25	3.75	14.0778		
	5	54.00	68.66	-14.66	215.0578		
	9	81.00	74.33	6.67	44.4566		
	13	93.00	80.00	13.00	169.0000		
	11	76.00	77.17	-1.17	1.3601		
	23	77.00	94.17	-17.17	294.7724		
	23	95.00	94.17	0.83	0.6907		
	28	109.00	101.25	7.75	60.0097		
intercept		61.5804		0.00	800.4305	= sum(res ²)	2274.00 SS total, from above
slope		1.4169					800.43 SS residual
							1473.57 SS improvement

4. What is the evidence? Calculate likelihood ratio for the overall model

The overall (omnibus) model includes all of the terms in the model.

The analytic model isolates a single term for comparison.

The omnibus and analytic models are the same for single variable models.

To calculate the likelihood ratio, we take the ratio of the total variance to the unexplained variance, then raise it to a power of sample size/2.

This can be calculated directly.

$$LR = (2274/800.4)^{9/2} = 110$$

This can also be calculated from the explained variance r^2 .

$$r^2 = SS_{\text{explained}} / SS_{\text{total}}$$

$$1 - r^2 = (SS_{\text{res}} / SS_{\text{total}})$$

The likelihood ratio is then calculated from the inverse of the unexplained variance.

$$LR = (1 - r^2)^{-n/2}$$

$$LR = (SS_{\text{res}} / SS_{\text{total}})^{-n/2}$$

$$LR = (800.4 / 2274)^{-9/2} = 110$$

The alternative model is 110 times more likely than the full (null) model.

We have good evidence in favor of the alternative over the null model.

5. Mode of inference.


This is an experiment under controlled conditions with a measurement protocol that could be repeated many times. We expect the results to vary from run to run due to many sources of variation including measurement error. From the law of large numbers (Bernoulli 1713) we expect better and better estimates of the true value of the parameters as we increase the number of repeats. We will use frequentist inference.

5. State population and whether sample is representative.

The population is a very large number of repeats of the experiment. We assume that our run of the experiment is representative of the population, for which the model is:

$$P_{\text{corn}} = \alpha + \beta_{P_{\text{Soil}}} \cdot P_{\text{Soil}} + \varepsilon$$

Running the experiment thousands of times, while conceivable, is hardly feasible. We will use the law of large number to infer from the measurements we have (the sample) to the larger population (thousands of runs). We will infer from our sample to the true value of the relation of phosphorus in corn to phosphorus in soil, as represented by the parameters $\beta_{P_{\text{Soil}}}$ and α . This view of inference emphasizes the importance of the experimental protocol. From this point of view, the population is all possible measurements, given the experimental protocol.



5. State population and whether sample is representative.

The population is not an enumerable population. It is not all corn plants in the world nor is it all corn plants in Iowa, where the experiment was conducted. It is not even all corn plants at the agricultural station where the experiment was conducted. If it were an enumerable population, which it could be, we would need to justify the sampling protocol as taking a representative sample from the enumerable population.

Looking beyond the results from a single run at a single location, can we infer beyond the location at the experimental agriculture station? Can we infer beyond the station location? Can we make an inference about the relation between phosphorus in corn and in soil in the state of Iowa? To do this, we need more information. For example, would the be independent of the 17 soils types in Iowa ? If we thought so, we would need to investigate this dependence by running the experiment across a range of soil types.

5. Decide on mode of inference. Is hypothesis testing appropriate?

If yes step 6, otherwise, calculate and report the likelihood ratio with parameter estimates.

Based on the law of large numbers we use frequentist inference. This allows us to infer to multiple runs of the experiment. With frequentist inference we use the likelihood ratio (a measure of strength of evidence) to calculate the uncertainty on this measure. Equivalently we ask: what is the cost (or risk) of drawing a false conclusion? In an applied context, such as this, one cost is recommending an agricultural practice, such as adding phosphorus to soil. In particular we wish to avoid recommending a practice with little or no effect. In such a context, we control our false conclusion rate at a fixed value. By convention this is 5%.

6. State H_A / H_0 , test statistic, and α

The term of interest in the model is $\beta_{Psoil} \cdot Psoil$.

The research hypothesis for this term is that phosphorus in corn depends on phosphorus in the soil, and hence variation in $\beta_{Psoil} \cdot Psoil$ term. With hypothesis testing the goal is to reject the null hypothesis, that of no relation and hence zero slope. $H_0: \beta_{Psoil} = 0$.

$$H_A: \beta_{Psoil} \neq 0$$

$$H_0: \beta_{Psoil} = 0$$

Hypotheses about terms in the model are equivalent to hypotheses about the parameters in the model.

$$H_A: \text{var}(\beta_{Psoil} \cdot Psoil) > 0 \quad \text{Equivalent to} \quad H_A: \beta_{Psoil} \neq 0$$

$$H_0: \text{var}(\beta_{Psoil} \cdot Psoil) = 0 \quad \text{Equivalent to} \quad H_0: \beta_{Psoil} = 0$$

6. State H_A / H_0 , test statistic, and tolerance for Type I error, α

The test statistic will be the F-statistic. This is a variance ratio, which we will construct in the next step.

The distribution will be the F-distribution, which is readily obtained from a cumulative frequency distribution..

We judged the errors to be acceptably homogeneous and normal, so later we will obtain a randomized p-value, free of assumptions, to check our judgement.

Tolerance for Type I error is $\alpha = 5\%$

7. ANOVA. Table Source, SS, df, MS, F-ratio.

Establish relation of model to ANOVA table.
GLM on left side of chalk board, with df and SS
ANOVA table headings at top, to right.
Fill in below GLM, then move Source, df, SS to table.
Move Source, df, and SS to ANOVA table
Complete calculations of MS, F,

Move Source, df, and SS to table. An ANOVA table is a GLM turned on its side.

Source	df	SS	MS	F	----> p
PSoil	1	1473.6			
<u>Res</u>	<u>7</u>	<u>800.4</u>			
Total	8	2274.0			

Complete the ANOVA table

$$MS = SS/df$$

$$MS_{\text{model}} = 1473.6/1 = 1473.6$$

$$MS_{\text{res}} = 800.4/7 = 114.34$$

$$F = MS_{\text{model}} / MS_{\text{res}} = 12.89$$

Computational flow is left to right,
compute MS from SS and df in ANOVA table
compute F from MS

Source	df	SS	MS	F	----> p
PSoil	1	1473.6	1473.6	12.89	
<u>Res</u>	<u>7</u>	<u>800.4</u>	114.34		
Total	8	2274.0			

Here are some additional statistics.

$$r^2 = \text{explained variance} = SS_{\text{model}} / SS_{\text{tot}}$$

r is the correlation coefficient

$$SS_{\text{model}} / SS_{\text{tot}} = \text{coefficient of determination, for any GLM}$$

$$1 - \text{coefficient of determination} = \text{coefficient of non-determination.}$$

7. Constructing the ANOVA table.

Here is a summary, before moving on to the next step in the recipe.

The ANOVA table represents a sequence of computations from left to right.

Factors are listed, based on the model that was written.

df are listed, for each factor, and for the residuals. $df_{\text{tot}} = df_{\text{model}} + df_{\text{res}}$

SS_{tot} is estimated as the sum of the squared deviations of the response variable from the grand mean of the response variable. This is partitioned into a component for each factor in the model, and one residual component.


MS means squares are computed for each source (model factors + residuals).

F ratios are formed as ratios of mean squares MS

7. ANOVA: Calculate Type I error from F distribution.

Packages compute and place the p-value in the ANOVA table. $p = 0.00885$

```
MTB > cdf 12.89 k3;
SUBC> f 1 8.
MTB > print k3
K3      0.991146
MTB > let k3 = 1-k3
MTB > print k3
K3      0.00885403
```



8. Recompute p-value to check our judgement concerning the residuals .

Confidence limits and p-values are computed from statistical distributions (χ^2 , F , t , and normal). However, these calculations can be inaccurate if assumptions for these distributions are violated. The distortion depends on sample size. As a rule of thumb distortion can be serious if $n < 30$, less serious if $30 \leq n \leq 100$, and usually not serious if $n > 100$. In this example $n = 9$, so violations are potentially serious.

When assumptions are not met, a recomputed Type I error is worth considering if:

n small (Yes, $n = 9$)

p near α (No, $p = 0.00885$)

The p-value is far from α and hence recomputing the p-value can easily change our estimate of Type I error, but will not change our decision. Randomized p-values generally differ from the theoretical p-values by less than a factor of two, rarely by a factor of 5 or more. This choice is a matter of judgement, through experience with randomized p-values. In this case, a factor of 5 would leave the p-value less than 5%, and hence not alter the decision. So we might well judge at this point not to recompute the p-value.

8. Recompute p-value to check our judgement concerning the residuals .

But while good judgement will suffice in day to day practice, it is not adequate when we must defend our conclusion from critical scrutiny. If we had to defend our conclusion in court, then we would need to be prepared to defend our result by use a p-value free of assumptions rather than appealing to judgement.

For the example at hand, it is of interest to find out whether our judgement was correct, that assumptions were met using the F-distribution to calculate a p-value.

In 4000 randomizations there were 27 instances of an F-ratio greater than 12.89.

The randomized p-value is somewhat smaller than the p-value from the F-distribution. The p-value from the F- distribution was high by a factor of

$$0.008854 / 0.00675 = 1.3$$

In 10,000 randomizations there were 79 instances of $F > 12.89$.

$$0.008854 / 0.0079 = 1.1$$

MTB > let k2 = 27/4000
MTB > print k2
K2 0.00675000

This is consistent with experience with randomization, which shows that the recalculated p-value will generally differ from that from the F-distribution by less than a factor of 2 (and rarely by more than a factor of 5).

Having calculated a randomized p-value, we will report it. The more randomizations the better, so we will report the empirical value from 10,000 randomization: $p = 0.0079$.

In the future, as our experience with judging assumptions from graphical displays grows we will have increased confidence in our judgement.

9. Declare and report decision about model terms.

$p = 0.0079$ (via randomization, hence no assumptions about error distribution)

$p < \alpha = 5\%$ so reject $H_0 : \text{var}(\beta_{P_{\text{Soil}}} \cdot P_{\text{Soil}}) = 0$

Rejecting the null on the variance is equivalent to rejecting the null on the parameter, $\beta_{P_{\text{Soil}}} \neq 0$

9. Report statistical conclusion.

We reject the no-effect hypothesis in favour of statistically significant increase in available phosphorus with increase in soil phosphorus

We report the test statistic, the sample size (or equivalent) and the p-value ($F_{1,7} = 12.89$, $p = 0.0079$ by randomization)

10. Report science conclusion. Interpret parameters of biological interest.

In this example our interest was in whether or not phosphorus content in corn was related to phosphorus content in the soil. However, we can learn more from the analysis than just this yes/no decision.

We begin by interpreting the slope parameter, estimated at was $\hat{\beta}_{P_{Soil}} = 1.4$ ppm/ppm.

The p-value from the F-distribution was close to that by randomization. The deviation from normality in the residual plot did not have a large effect on the p-value estimate. So we can be confident that the deviation in the residual plot did not have a large effect on the parameter estimate.

The observed phosphorus content in corn increases by 1.4 units for each unit increase in soil phosphorus. Are the results consistent with amplification in corn relative to the soil? In other words, can we exclude the hypothesis ($\beta_{P_{Soil}} = 1$) that there was no amplification?

To avoid repeated testing, we compute the confidence limits using the standard error of the estimate of the slope parameter $s_b = 0.3947$. This estimate is provided by the statistical package. It differs from the standard error of the mean.

95% confidence limits for parameters.

$$P\{\text{Lower} < \beta_{P_{Soil}} < \text{Upper}\} = 1 - \alpha = 95\%$$

$$\text{Lower} = \hat{\beta}_{P_{Soil}} - t_{0.025[7]} * s_b$$

$$\text{Lower} = 1.4169 - 2.3646 * 0.3947 = 0.484$$

$$\text{Upper} = \hat{\beta}_{P_{Soil}} + t_{0.025[7]} * s_b$$

$$\text{Upper} = 1.4169 + 2.3646 * 0.3947 = 2.35$$

$$0.484 \leq \beta_{P_{Soil}} \leq 2.35 \text{ for } \alpha = 5\%.$$

These limits exclude the $\beta_{P_{Soil}} = 0$ hypothesis. They do not exclude the hypothesis $\beta_{P_{Soil}} = 1$ hypothesis, that available phosphorus increases in direct proportion to soil phosphorus. We cannot reject the hypothesis that phosphorus in corn increases in direction proportion to phosphorus in soil $\beta_{P_{Soil}} = 1$.

Confidence limits allow us to exclude several hypotheses, not just the null of no relation.

For the sake of completeness, we report the regression equation with standard error of the slope parameter and sample size.

$$P_{corn} = 61.58 + 1.42 P_{Soil} \quad s_b = 0.3947, n = 9$$

References

Bernoulli, Jakob (2005) [1713], *The Art of Conjecturing*, together with Letter to a Friend on Sets in Court Tennis (English translation), translated by Edith Sylla. Baltimore: Johns Hopkins Univ Press.

Hurlbert, S.H. (1984) Pseudoreplication and the design of ecological field experiments. *Ecological Monographs* 54: 187-211

Snedecor, G.W. and Cochran, W.G. (1989) *Statistical Methods*. 8th Edition. Ames: Iowa State University Press.